Flight Delay Analysis: Understanding How Delays Actually Behave

Purpose	Document
Style	Analytical
Version	4

📌 The Problem — Are Delays Really Random?

Flight delays are commonly treated as unpredictable disruptions caused by external factors such as weather, congestion, or operational inefficiencies. Because of this, delays are often assumed to be random and difficult to manage proactively.

However, this assumption creates a major limitation. If delays are viewed as random, organizations are less likely to identify repeatable patterns or implement systematic improvements.

This analysis challenges that assumption by asking:

> Are flight delays truly random, or do they follow consistent patterns that can be measured and improved?

🎯 Objective

The objective of this analysis is to:

- Identify the key drivers of flight delays
- Understand how delays evolve across time and operations
- Validate whether delay behavior is random or systematic

🧠 Analytical Approach

Rather than assuming a single cause, this analysis breaks the problem into multiple dimensions and tests each independently:

- Time-based behavior (when delays occur)
- Volume vs delay relationship (does traffic drive delays?)
- Distribution of delays (how delays are spread)
- Airport-level performance differences

The approach follows a structured process:

> Identify patterns → Test relationships → Validate with statistics → Translate into operational insights

---

📊 Data & Methodology

---

4. The Dashboard: Designed for Decisions, Not Just Data

This analysis was supported by a 2-page Power BI dashboard designed to prioritize clarity and decision-making speed.

Design Principle:

> Executives and stakeholders should understand the story within seconds.

Key Features:

- KPI overview (Total Flights, Avg Delay, % Delayed)
- Time-based trends (hourly delay patterns)
- Volume vs delay relationship (scatter analysis)
- Airport-level performance comparison
- Distribution of delays (severity breakdown)

Dashboard Overview

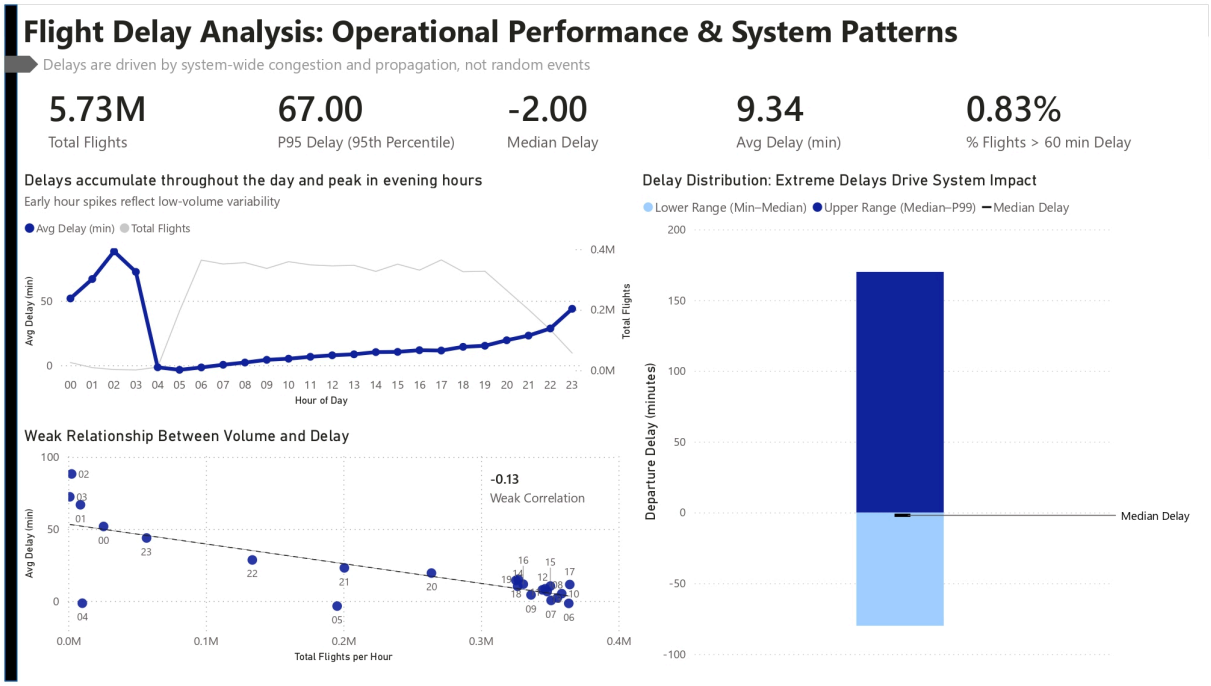

Figure 1: Overview dashboard highlighting delay distribution, time-based behavior, and volume relationship.

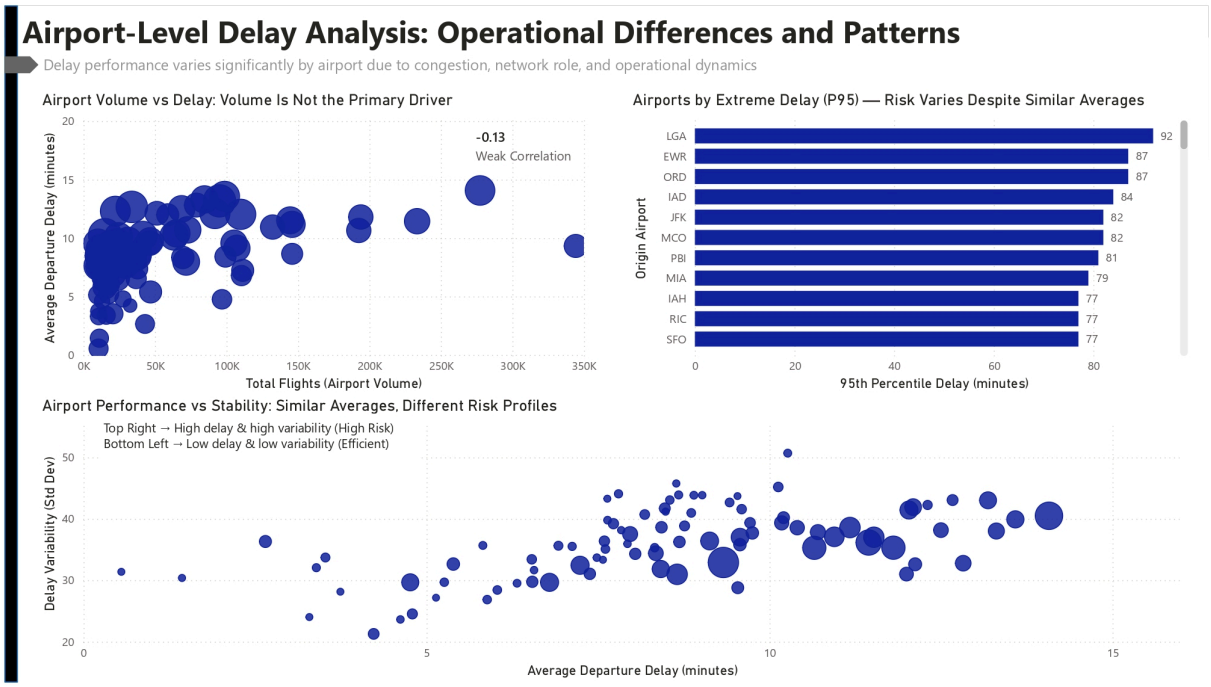

Figure 2: Airport-level analysis showing performance differences using extreme delay metrics (P95), variability, and volume.

📊 Data & Methodology

Data Source

- Dataset: *2015 Flight Delays and Cancellations (Kaggle)*
- Total records: ~5.8 million flights
- Filtered dataset: ~5.7 million flights (after cleaning)

Key Metric

- Departure Delay (minutes)

Data Preparation

- Removed NULL values in departure delay
- Excluded cancelled flights from delay calculations
- Converted time into hourly buckets (`dep_hour`)

Time Segmentation

- Early Hours: 00:00–05:59
- Core Operations: 06:00–17:59
- Evening Peak: 18:00–23:59

🔍 Analysis & Statistical Validation

0. Airport-Level Performance Analysis

To assess whether delay behavior varies by location, airport-level performance was analyzed using multiple metrics, including average delay, variability, extreme delays (P95), and proportion of severe delays.

Approach:

- Aggregated departure delays by origin airport
- Applied a minimum volume threshold (≥10,000 flights) to ensure statistical reliability
- Evaluated performance using:
  - Average delay (overall efficiency)
  - Standard deviation (operational stability)
  - 95th percentile delay (extreme disruption)
  - Severe delay ratio (>60 minutes)

Key Observations (High-Volume Airports):

- **Volume does not determine delay performance:**

  High-volume airports such as ORD (~277K flights) show only slightly higher delays compared to lower-volume airports like DAL (~59K flights), reinforcing that traffic volume alone does not drive delays.

- **Airports with similar averages show different risk profiles:**

  For example, LGA and DAL have similar average delays (~13 vs ~12 minutes), but LGA has significantly higher variability (std dev ~43 vs ~31) and higher extreme delays (P95 ~92 vs ~65), indicating greater operational instability.

- **Operational stability varies significantly:**

  DAL exhibits the lowest variability and lowest severe delay ratio (~5.4%), suggesting more controlled and predictable operations, while LGA shows the highest variability and severe delay ratio (~8.4%), indicating higher susceptibility to disruption.

- **High volume does not imply instability:**

  ORD, despite being the busiest airport in the dataset, maintains moderate delay metrics (avg ~14, P95 ~88), demonstrating that efficient operations can offset the impact of high traffic.

Interpretation:

Airport performance differences are driven by **operational efficiency, stability, and ability to manage disruptions**, rather than traffic volume alone. Variability and

extreme delays provide a more accurate representation of operational risk than averages.

Conclusion:

> Airport-level analysis shows that delay behavior varies significantly across locations due to operational factors. High-volume airports are not necessarily less efficient, and performance differences are better explained by variability and extreme delay behavior rather than average delay alone.

---

1. Volume vs Delay Relationship

To test whether traffic volume drives delays, correlation analysis was performed at a granular (airport-hour) level.

- Correlation coefficient: -**0.13**

Interpretation:

The relationship between flight volume and delay is negligible. The slight negative correlation suggests that higher-volume periods may operate more efficiently due to better coordination and resource allocation.

Conclusion:

> Flight volume alone does not explain delay behavior.

---

2. Time-Based Delay Behavior

Segment	Avg Delay (min)	Std Dev	Sample Size
Early Hours	6	46.9	~244K
Core Operations	6	31.2	~4.18M
Evening Peak	19	48.6	~1.31M

Insights:

- Delay levels remain stable (~6 minutes) during early and core operating hours
- Delays increase sharply to ~19 minutes in the evening
- Variability is highest in the evening, indicating reduced operational stability

Interpretation:

Delays do not build gradually throughout the day. Instead, operations remain relatively stable until disruptions accumulate and reach a tipping point in later hours.

Additional Note:

Early-hour delay spikes are influenced by low flight volume, which increases variability and makes average delay more sensitive to individual extreme events. As a result, early-hour values may appear disproportionately high despite limited operational impact.

Conclusion:

> Delay behavior is time-dependent and reflects accumulated operational strain rather than isolated events.

---

3. Delay Distribution (Extreme Value Analysis)

Metric	Delay (min)
Median	-2
P90	35
P95	67
P99	168
Max	1988

Note on Metric Selection:

The 95th percentile (P95) is used as the primary indicator of extreme delay behavior. It provides a stable measure of high-impact delays while reducing sensitivity to extreme outliers compared to maximum values.

Insights:

- Median delay is negative → most flights depart slightly early
- Delays increase sharply in upper percentiles
- Extreme delays reach nearly 2000 minutes

Interpretation:

The distribution is heavily right-skewed, meaning that while most flights operate close to schedule, a small percentage experience severe delays.

Conclusion:

> A small number of extreme delays disproportionately impact overall system performance.

🔍 Key Findings: What the Data Reveals

🔍 Finding 1: Delays Follow a Time-Based Pattern

- Delays remain stable during early and core hours
- Increase sharply during evening peak

Implication:

➡️ Delays accumulate rather than occur randomly

➡️ Early disruptions propagate throughout the day

Takeaway:

👉 Improving early-day operations can reduce system-wide delays

---

🔍 Finding 2: Flight Volume Does Not Drive Delays

- Correlation between volume and delay: -0.13
- High-volume airports do not consistently show higher delays

Implication:

➡️ Traffic volume is not the primary driver

➡️ Operational efficiency plays a larger role

Takeaway:

👉 Reducing traffic alone will not solve delay issues

🔍 Finding 3: Delays Are Driven by Extreme Events

- Median delay: -2 minutes (most flights on time/early)
- P99 delay: 168 minutes
- Max delay: 1988 minutes
- P95 is used as the primary indicator of extreme delay behavior due to its stability and resistance to extreme outliers compared to maximum values.

Implication:

➡️ A small number of extreme delays dominate system performance

Takeaway:

👉 Managing high-impact delays is more effective than optimizing averages

---

🔍 Finding 4: Airport Performance Varies Significantly

- Airports with similar volume show different delay behavior
- Variability and extreme delays differ across locations
- These differences are further visualized through comparative analysis of average delay versus variability, highlighting how similar average performance can mask underlying instability.

Implication:

➡️ Performance differences are operational, not volume-driven

➡️ Some airports are more stable, others more disruption-prone

Takeaway:

👉 Improvement strategies should be location-specific

🧠 Key Findings

1. **Delays are not random** — they follow consistent and observable patterns
2. **Flight volume is not a primary driver** — correlation is negligible
3. **Delays escalate in the evening** — indicating accumulation of disruptions
4. **Extreme delays drive system performance** — distribution is heavily skewed

---

🔗 System-Level Insight

Combining all findings reveals a consistent pattern:

- Early disruptions are not fully absorbed
- Effects propagate throughout the day
- System performance is driven by high-impact events rather than average conditions

This can be understood as a **connected system** where delays interact and compound over time.

> Delays are not isolated incidents — they are system-level outcomes driven by operational flow.

🔧 Recommendations

1. Stabilize Early-Day Operations

- Reduce delays at the start of the day
- Improve readiness of aircraft and crew
- Prevent downstream propagation of disruptions

2. Focus on High-Impact Delays

- Prioritize flights with delays >60 minutes
- Allocate resources dynamically to high-risk cases
- Mitigate cascading effects early

3. Optimize Evening Operations

- Address accumulation effects
- Adjust scheduling to reduce late-day congestion
- Improve recovery strategies

4. Implement Real-Time Monitoring

- Track delay buildup throughout the day
- Detect early warning signals
- Enable proactive intervention

⚠️ Limitations & Considerations

While this analysis provides strong evidence of structured delay behavior, several limitations should be considered:

1. Use of Departure Delay Only

- Analysis focuses on departure delays
- Does not fully capture downstream effects (e.g., arrival delays, network propagation)

👉 Impact: May underestimate system-wide delay interactions

---

2. Limited External Factors

- Weather, airline-specific operations, and aircraft constraints were not included

👉 Impact: Some delay drivers remain unobserved

---

3. Single-Year Dataset (2015)

- Results are based on one year of data

👉 Impact: Patterns may vary across different years or macro conditions

---

4. Aggregation Effects

- Some analyses rely on aggregated metrics (hourly, airport-level)

👉 Impact: May smooth out micro-level variability

---

5. Volume Filtering Assumptions

- Airports with low flight volume were excluded to improve statistical reliability

👉 Impact: Smaller airports and niche patterns are not represented

🎯 Final Takeaway

> Flight delays are driven by operational dynamics and extreme events, not by average conditions or traffic volume.

Effective delay management requires:

- Early intervention
- Focus on high-impact disruptions
- System-level optimization

🛠️ Tools Used

- SQL Server (data extraction, transformation, and analysis)
- Power BI (dashboarding and visualization)

📌 Closing Statement

This analysis demonstrates that understanding how delays behave — rather than simply measuring them — enables more effective and targeted operational improvements.

By shifting from averages to system behavior, organizations can better control how delays develop and prevent them from impacting overall performance.